

# Apache Hive: A Petabyte-Scale Data Warehouse System Over a MapReduce Framework

Test Application: MBV Climate and Ocean Intelligence Africa

Dushime Mudahera Richard

Databases for Big Data Seminar – Prof. Iztok Savnik

January 2026

# Agenda & Executive Summary

## Outline

- ① System Architecture
- ② Query Optimization
- ③ Experimental Results
- ④ Challenges & Solutions
- ⑤ Conclusions

## Project Overview

Production-grade **distributed big data ecosystem** simulating climate analytics

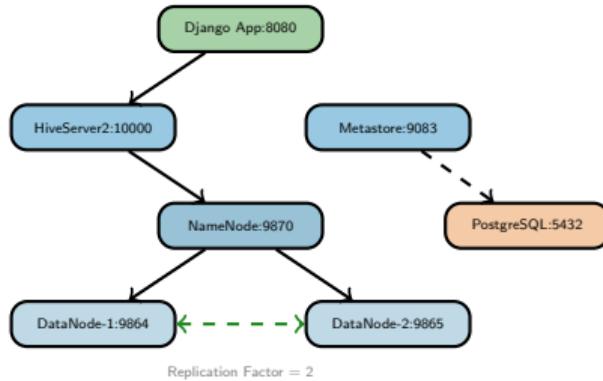
-  7-Container Docker Stack – Hadoop 3 + Hive 2.3.2
-  Django REST API – Climate dashboard & benchmarking
-  3-Node HDFS Cluster – Distributed storage with replication

## Test Data

- 4.75 million climate records
- 304 MB raw CSV data
- 4 Hive tables (1980–2024)
- 5 African regions, 44 countries

**Goal:** Demonstrate Apache Hive's capability for petabyte-scale analytics on commodity hardware

# 7-Container Stack Architecture



## Container Services

- **Django** – REST API, Web Dashboard
- **HiveServer2** – JDBC/Thrift gateway
- **Metastore** – Schema catalog (decoupled)
- **NameNode** – HDFS namespace manager
- **DataNodes** – Distributed block storage
- **PostgreSQL** – Persistent metadata DB

## HDFS Cluster Stats

- Configured Capacity: 416.56 GB
- Live DataNodes: 2 (healthy)
- Under-replicated blocks: 0

# Apache Hive & Query Optimization

## Hive Architecture

- **HiveServer2** – JDBC/ODBC gateway
- **Metastore** – Centralized schema catalog
- **Execution Engine** – MapReduce backend
- **HDFS Storage** – Distributed file system

## Cost-Based Optimizer (CBO)

- Uses Apache Calcite for query planning
- Estimates costs from table statistics
- Selects optimal join algorithms
- Requires ANALYZE TABLE

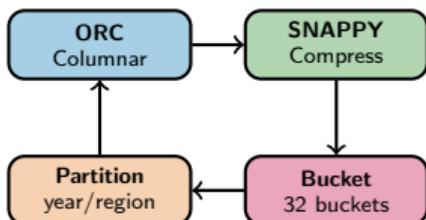
## Join Algorithms

- **Shuffle Join (Reduce-Side)**
  - Default, requires data shuffle
  - High network I/O cost
- **Broadcast Join (Map-Side)**
  - Small table broadcast to all mappers
  - 3x faster for asymmetric joins
- **Sort-Merge-Bucket Join**
  - Most efficient for pre-bucketed data

## Key Configuration

```
SET hive.auto.convert.join=true;
```

# Storage Optimizations & Performance



## Optimization Techniques

- **ORC Format** – Columnar storage, predicate pushdown, reduced I/O
- **SNAPPY** – Fast compression, CPU-efficient
- **Partitioning** – Skip 90%+ data via partition pruning
- **Bucketing** – Enable Map-Side joins, faster GROUP BY

## Compression Results

- ORC: **88%** size reduction
- vs. raw CSV (304 MB → 36 MB)

## Vectorized Execution

- Process 1,024 rows per CPU instruction
- `SET hive.vectorized.execution.enabled=true;`
- Significant speedup for STDDEV, CORR, AVG

# Query Performance Benchmarks

## Query Execution Times (4.75M rows)

Query Type	Time (s)
Simple Regional Aggregation	<b>9.31</b>
Complex Monthly Aggregation	14.88
Statistical Analysis	16.13
Yearly Analysis	15.24
Map-Side Join	15.30
Reduce-Side Join	42.70
Data Coverage Query	13.37

## Regional Temperature Results

Region	Avg °C	Variance
Central	30.49	2.448
West	29.51	2.448
East	26.51	2.448
North	24.51	2.450
South	20.50	2.448

## Join Algorithm Comparison

- Map-Side: **15.30s** (2.8x faster)
- Reduce-Side: 42.70s (baseline)

## Key Findings

- 10°C gradient Central→South
- Uniform variance (~ 2.45)
- 2-stage MapReduce for ORDER BY
- 28–35 GB HDFS reads per query

# Engineering Challenges & Solutions

## Challenge 1: ARM Emulation

- Apple Silicon runs x86 images
- 10–20% performance overhead
- 60–120s container startup

## Challenge 2: PostgreSQL Driver

- Hive 2.3.2 JDBC compatibility
- MD5 auth required (PG 9.6)
- Metastore schema init

## Solution

- Extended health check timeouts
- Optimized JVM heap settings
- Rosetta 2 emulation layer

## Solution

- Manual JAR injection
- postgresql-42.7.2.jar
- Docker volume mounts

# Summary & Key Takeaways

## Achievements

### • Infrastructure

- 7-container Docker stack
- Hive 2.3.2 + Hadoop 3
- 447 GB HDFS capacity

### • Data Processing

- 4.75M records (45 years)
- 44 countries, 5,000 stations
- Multi-stage MapReduce

### • Performance

- 2.8x speedup with Map-Side joins
- 88% compression with ORC
- 9–26s query latency

## Key Findings

- ➊  CBO requires fresh statistics
- ➋  Join algorithm choice is critical
- ➌  ORC format essential for analytics

## Conclusion

Apache Hive provides cost-effective, petabyte-scale analytics on commodity hardware—validated by Netflix, Facebook, Airbnb processing 100+ PB daily.

# Thank You!

Questions?



**Dushime Mudahera Richard**

*MBV Climate and Ocean Intelligence Africa*

Databases for Big Data Seminar – Prof. Iztok Savnik

University of Primorska — January 2026